

# Machine Learning

## Lecture 3: Training Models

**Asst. Prof. Dr. Santitham Prom-on**

Department of Computer Engineering, Faculty of Engineering  
King Mongkut's University of Technology Thonburi

# Topics

- Direct approach: linear regression with ordinary least square
- Iterative approach
  - Gradient descent
  - Batch gradient descent
  - Stochastic gradient descent
  - Mini-batch gradient descent
- Logistic regression

# Model training

Two different ways

- Using a direct “closed-form” equation.
- Using an iterative optimization approach.

# Direct approach

## Linear regression with Ordinary Least Square (OLS)

- Linear regression model can be expressed by the following formula

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n + \varepsilon$$

where

$\hat{y}$  is the predicted value

n is the number of features

$x_i$  is the  $i^{\text{th}}$  feature value

# Simple linear regression

$$Y_i = \alpha + \beta X_i + \varepsilon_i$$

$$e_i \sim N(0, \sigma^2) \quad i.i.d.$$

$\varepsilon_i$  is independent of  $X_i$

- The intercept is  $\alpha$
- The slope is  $\beta$
- We use the normal distribution to describe the “error”


# Method of least squares

- Choose the  $\beta$ 's so that the sum of the squares of the errors,  $\varepsilon_i$ , are minimized
- The least squares function is

$$\begin{aligned} S &= \sum_{i=1}^n \varepsilon_i^2 \\ &= \sum_{i=1}^n (y_i - \alpha - \beta x_i)^2 \end{aligned}$$

# OLS solution

Minimum of a function is the point where the slope is zero



# Derivative of the error functions

The function  $S$  is to be minimized with respect to  $\beta_0, \beta_1$

and

$$\frac{\partial S}{\partial \alpha} = -2 \sum_{i=1}^n (y_i - \alpha - \beta x_i) = 0$$

$$\frac{\partial S}{\partial \beta} = -2 \sum_{i=1}^n (y_i - \alpha - \beta x_i) x_i = 0$$



# Least square normal equation

$$n\alpha + \beta \sum_{i=1}^n x_i = \sum_{i=1}^n y_i$$

$$\alpha \sum_{i=1}^n x_i + \beta \sum_{i=1}^n x_i^2 = \sum_{i=1}^n x_i y_i$$

# Find alpha (intercept)

$$\alpha = \frac{\begin{vmatrix} \sum_{i=1}^n y_i & \sum_{i=1}^n x_i \\ \sum_{i=1}^n x_i y_i & \sum_{i=1}^n x_i^2 \end{vmatrix}}{\begin{vmatrix} n & \sum_{i=1}^n x_i \\ \sum_{i=1}^n x_i & \sum_{i=1}^n x_i^2 \end{vmatrix}} = \frac{\sum_{i=1}^n x_i^2 \sum_{i=1}^n y_i - \sum_{i=1}^n x_i y_i \sum_{i=1}^n x_i}{n \sum_{i=1}^n x_i^2 - \left( \sum_{i=1}^n x_i \right)^2}$$

# Find beta (slope)

$$\beta = \frac{\begin{vmatrix} n & \sum_{i=1}^n y_i \\ \sum_{i=1}^n x_i & \sum_{i=1}^n x_i y_i \end{vmatrix}}{\begin{vmatrix} n & \sum_{i=1}^n x_i \\ \sum_{i=1}^n x_i & \sum_{i=1}^n x_i^2 \end{vmatrix}}} = \frac{n \sum_{i=1}^n x_i y_i - \sum_{i=1}^n x_i \sum_{i=1}^n y_i}{n \sum_{i=1}^n x_i^2 - \left( \sum_{i=1}^n x_i \right)^2}$$

# Example: Multiple regression

$x_1$	$x_2$	$y$
1	2	12
2	1	9
3	2	19
1	1	8

$$y = c_1x_1 + c_2x_2 + c_3$$

$$\begin{aligned}c_1 + 2c_2 + c_3 &= 12 \\2c_1 + c_2 + c_3 &= 9 \\3c_1 + 2c_2 + c_3 &= 19 \\c_1 + c_2 + c_3 &= 8\end{aligned}$$

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 1 & 1 \\ 3 & 2 & 1 \\ 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} c_1 \\ c_2 \\ c_3 \end{bmatrix} = \begin{bmatrix} 12 \\ 9 \\ 19 \\ 8 \end{bmatrix}$$

# Pseudoinverse

**Y**  $N \times 1$  vector

**A**  $N \times M$  matrix, where  $M$  is the number of parameters

**B**  $M \times 1$  vector

$$\mathbf{Y} = \mathbf{AB}$$

$$\mathbf{AB} = \mathbf{Y}$$

$$\mathbf{A}^T \mathbf{AB} = \mathbf{A}^T \mathbf{Y}$$

$$\underbrace{(\mathbf{A}^T \mathbf{A})^{-1}} \mathbf{A}^T \mathbf{AB} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{Y}$$

$$\mathbf{B} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{Y}$$

# Python: pseudoinverse

```
A = np.array([[1,2,1],  
              [2,1,1],  
              [3,2,1],  
              [1,1,1]])  
  
print(A)
```

```
[[1 2 1]  
 [2 1 1]  
 [3 2 1]  
 [1 1 1]]
```

```
B = np.array([12,9,19,8])  
B.shape = (-1,1)  
print(B)
```

```
[[12]  
 [ 9]  
 [19]  
 [ 8]]
```

```
np.matmul(np.linalg.pinv(A),B)
```

```
array([[ 3. ],  
       [ 5.5],  
       [-1.5]])
```

# Iterative approach

## Gradient descent

- A very generic optimization algorithm capable of finding optimal solutions to a wide range of problems.




Figure 4-3. Gradient Descent

Suppose you are lost in the mountains in a dense fog; you can only feel the slope of the ground below your feet.

A good strategy to get to the bottom of the valley quickly is to go downhill in the direction of the steepest slope.

# Nonlinear Least Square

- Suppose that we have a sample of  $n$  observations on the response and the regressor, say,  $y_i, x_{i1}, x_{i2}, \dots, x_{ik}$  for  $i=1,2,\dots,n$
- The least square method involves minimizing the least square function


$$S(\boldsymbol{\beta}) = \sum_{i=1}^n [y_i - f(\mathbf{x}_i, \boldsymbol{\beta})]^2$$



# Nonlinear objective function

Non-linear  
objective function

Unsolvable for roots  
of derivative of error




$$\frac{dE}{dx} = e^{-x}$$

# Gradient Descent Basic 1


## Objective function

Minimum of a function is found by following the slope of the function




# Gradient Descent Basic 2

## Moving opposite to gradient




# Gradient Descent Basic 3

## Iterative gradient evaluation




# Gradient Descent Basic 4

## Moving opposite to gradient




# Gradient Descent Basic 5

Iteratively descent opposite to the gradient




# Gradient Descent – algorithm

- Start with a point (randomly guessing)
- Repeat
  - Determine a descent direction
  - Choose a step
  - Update
- Until stopping criterion is satisfied



# Gradient Descent – algorithm


- Start with a point (randomly guessing)
- Repeat
  - Determine a descent direction
  - Choose a step
  - Update
- Until stopping criterion is satisfied






# Gradient Descent – algorithm

- Start with a point (randomly guessing)
- Repeat
  - Determine a descent direction
  - Choose a step
  - Update
- Until stopping criterion is satisfied




# Gradient Descent – algorithm

- Start with a point (randomly guessing)
- Repeat
  - Determine a descent direction
  - Choose a step
  - **Update**
- Until stopping criterion is satisfied



# Gradient Descent – algorithm

- Start with a point (randomly guessing)
- Repeat
  - Determine a descent direction
  - Choose a step
  - Update
- Until stopping criterion is satisfied



# Gradient Descent – algorithm

- Start with a point (randomly guessing) → Randomly guessing  $\beta$
- Repeat
  - Determine a descent direction →  $\text{direction} = -\frac{dS(\beta)}{d\beta}$
  - Choose a step →  $\text{step} > 0$
  - Update →  $\beta^{t+1} = \beta^t - \text{step} \frac{dS(\beta)}{d\beta}$
- Until stopping criterion is satisfied →  $\frac{dS(\beta)}{d\beta} \approx 0$

# Batch Gradient Descent

- To implement Gradient Descent, you need to compute the gradient of the cost function with regards to each model parameter  $\theta_j$ .
- Batch gradient descent uses data of the whole batch to compute gradient

$$\frac{\partial}{\partial \theta_j} \text{MSE}(\boldsymbol{\theta}) = \frac{2}{m} \sum_{i=1}^m (\boldsymbol{\theta}^T \mathbf{x}^{(i)} - y^{(i)}) x_j^{(i)}$$

$$\boldsymbol{\theta}^{(\text{next step})} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} \text{MSE}(\boldsymbol{\theta})$$

# Stochastic gradient descent

- Batch Gradient Descent uses the whole training set to compute the gradients at every step, which makes it very slow when the training set is large.
- Stochastic gradient descent samples random instances for training at each training step




Figure 4-9. Stochastic Gradient Descent

# Mini-Batch Gradient Descent

- Both stochastic and use small batch
- Apply for small batch instead of an instance.
- Faster by GPU computing

---

**Algorithm 8.1** Stochastic gradient descent (SGD) update at training iteration  $k$

---

**Require:** Learning rate  $\epsilon_k$

**Require:** Initial parameter  $\theta$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

    Compute gradient estimate:  $\hat{\mathbf{g}} \leftarrow +\frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$ .


    Apply update:  $\theta \leftarrow \theta - \epsilon \hat{\mathbf{g}}$ .

**end while**

---

# Example: Logistic Regression

Classification





# Logistic regression

Logistic regression has the following mathematical formulae,

$$\left( \frac{H_{\theta}(x_i)}{1 - H_{\theta}(x_i)} \right) = f(x_i) = \theta_0 + \sum_{i=1}^n \theta_i x_i$$

and

$$H_{\theta}(x_i) = \frac{1}{1 + e^{-f(x)}}.$$

# Loss function

This function has a nice property that,

$$\frac{\partial}{\partial \theta} H_{\theta}(x_i) = H_{\theta}(x_i)(1 - H_{\theta}(x_i))x_i.$$

For logistic regression, we can formulate the loss function as follows,

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n [-y_i \log H_{\theta}(x_i) - (1 - y_i) \log (1 - H_{\theta}(x_i))].$$

# Gradient

$$\begin{aligned}
 \frac{\partial J(\boldsymbol{\theta})}{\partial \boldsymbol{\theta}} &= \frac{\partial}{\partial \boldsymbol{\theta}} \left( \frac{1}{n} \sum_{i=1}^n (-y_i \log H_{\boldsymbol{\theta}}(x_i) - (1 - y_i) \log (1 - H_{\boldsymbol{\theta}}(x_i))) \right) \\
 &= \frac{1}{n} \sum_{i=1}^n \left( -y_i \frac{\partial}{\partial \boldsymbol{\theta}} \log H_{\boldsymbol{\theta}}(x_i) - (1 - y_i) \frac{\partial}{\partial \boldsymbol{\theta}} \log (1 - H_{\boldsymbol{\theta}}(x_i)) \right) \\
 &= \frac{1}{n} \sum_{i=1}^n \left( -\frac{y_i}{H_{\boldsymbol{\theta}}(x_i)} \frac{\partial}{\partial \boldsymbol{\theta}} H_{\boldsymbol{\theta}}(x_i) - \frac{1 - y_i}{1 - H_{\boldsymbol{\theta}}(x_i)} \frac{\partial}{\partial \boldsymbol{\theta}} (1 - H_{\boldsymbol{\theta}}(x_i)) \right) \\
 &= \frac{1}{n} \sum_{i=1}^n \left( -\frac{y_i}{H_{\boldsymbol{\theta}}(x_i)} H_{\boldsymbol{\theta}}(x_i) (1 - H_{\boldsymbol{\theta}}(x_i)) x_i \right. \\
 &\quad \left. - \frac{1 - y_i}{1 - H_{\boldsymbol{\theta}}(x_i)} H_{\boldsymbol{\theta}}(x_i) (1 - H_{\boldsymbol{\theta}}(x_i)) (-x_i) \right) \\
 &= \frac{1}{n} \sum_{i=1}^n \left( -y_i (1 - H_{\boldsymbol{\theta}}(x_i)) x_i - (1 - y_i) H_{\boldsymbol{\theta}}(x_i) (-x_i) \right) \\
 &= \frac{1}{n} \sum_{i=1}^n \left( -y_i (1 - H_{\boldsymbol{\theta}}(x_i)) + (1 - y_i) H_{\boldsymbol{\theta}}(x_i) \right) x_i \\
 &= \frac{1}{n} \sum_{i=1}^n \left( -y_i + y_i H_{\boldsymbol{\theta}}(x_i) + H_{\boldsymbol{\theta}}(x_i) - y_i H_{\boldsymbol{\theta}}(x_i) \right) x_i \\
 &= \frac{1}{n} \sum_{i=1}^n (H_{\boldsymbol{\theta}}(x_i) - y_i) x_i
 \end{aligned}$$


# Gradient descent

We can then iteratively update the parameters using the gradient descent approach,

$$\begin{aligned}\boldsymbol{\theta}^{(k+1)} &= \boldsymbol{\theta}^{(k)} - \alpha \frac{\partial J\boldsymbol{\theta}}{\partial \boldsymbol{\theta}} \\ &= \boldsymbol{\theta}^{(k)} - \alpha \left( \sum_{i=1}^n (H_{\boldsymbol{\theta}^{(k)}}(x_i) - y_i) x_i \right),\end{aligned}$$

where  $\alpha$  is the learning rate. The initial parameters  $\boldsymbol{\theta}^{(0)}$  can be randomized or set to any values as the loss function is convex.

# Results



# End of Lecture 3

## Question?