


Reinforcement Learning

Overview

- ▣ What is Reinforcement Learning (RL)?
- ▣ Markov Decision Processes
- ▣ Q-Learning

Types of Learning



Source: cognub.com

Supervised Learning

- **Data:** (x,y)
 - x is data
 - y is label
- **Goal:** Learn a function to map $x \rightarrow y$
- **Example:** Classification, regression, object detection, semantic segmentation, image captioning, etc.




CLASSIFICATION

→ cat

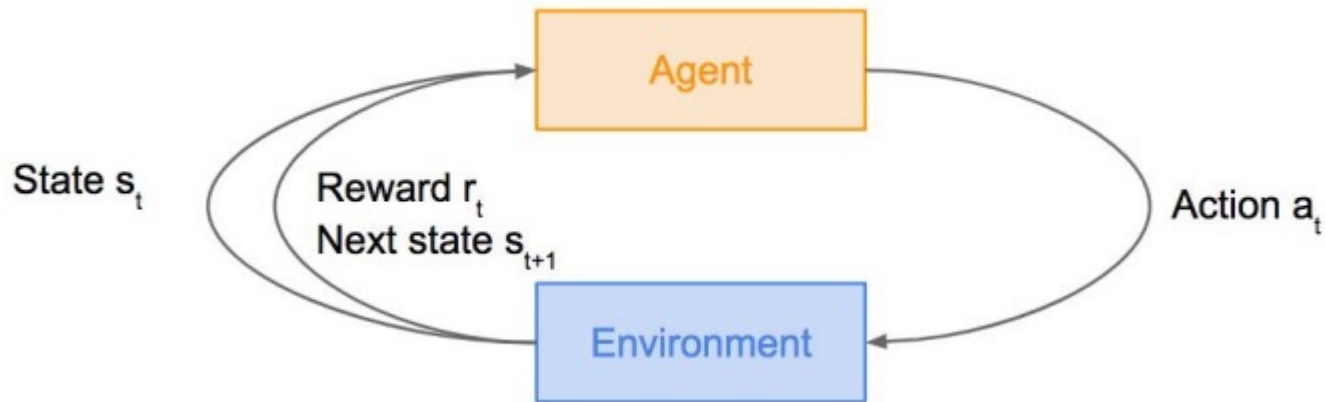
Unsupervised Learning

- **Data:** x
 - Just data, NO label!
- **Goal:** Learn some underlying *hidden structure* of the data
- **Example:** Clustering, dimensionality reduction, feature learning, density estimation, etc.



Reinforcement Learning

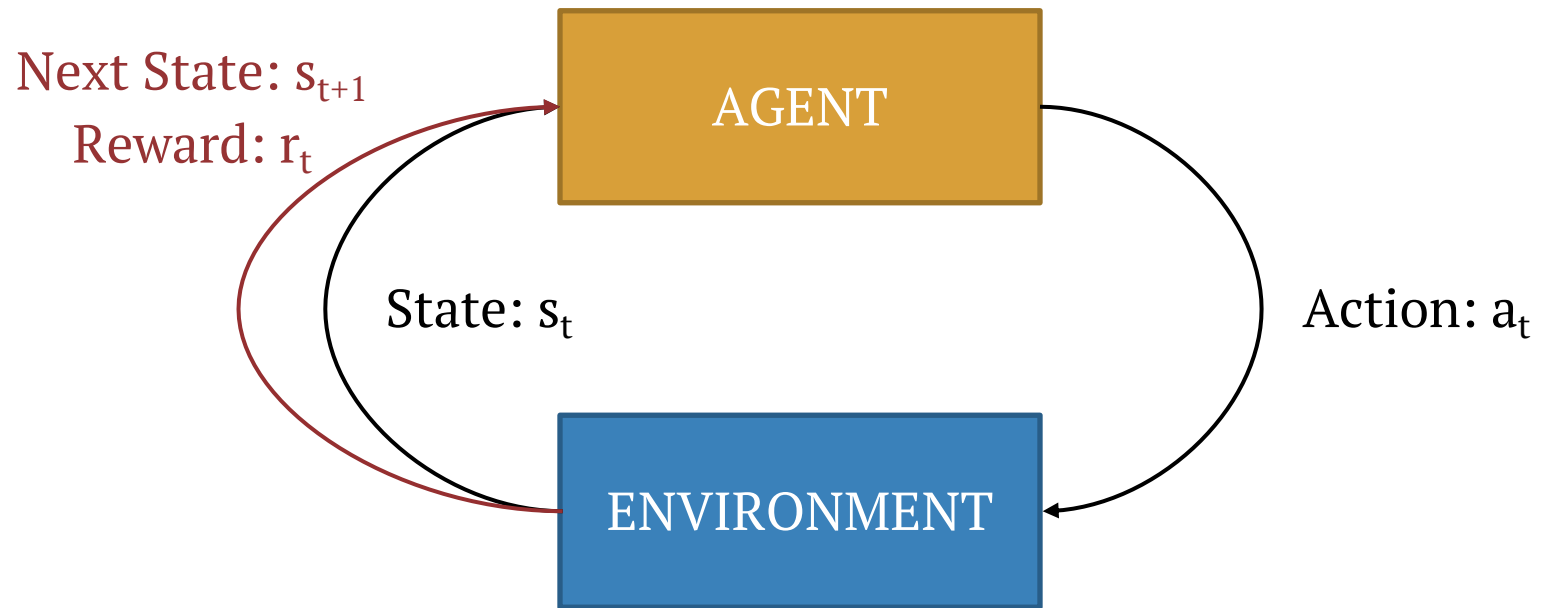
- ▣ Problems involving an **AGENT** interacting with an **ENVIRONMENT**, which provides numeric **REWARD** signals
- ▣ Goal: Learn how to take actions in order to maximize reward



Atari's Arcade Game



Reinforcement Learning



Atari Games


- ▣ **Objective:** Complete the game with the highest score



- ▣ **State:** Raw pixel inputs of the game state
- ▣ **Action:** Game controls e.g. Left, Right, Up, Down
- ▣ **Reward:** Score increase/decrease at each time step

Robot Locomotion

- ▣ **Objective:** Make the robot move forward
- ▣ **State:**
- ▣ **Action:**
- ▣ **Reward:**




Go

■ **Objective:** Win the game!

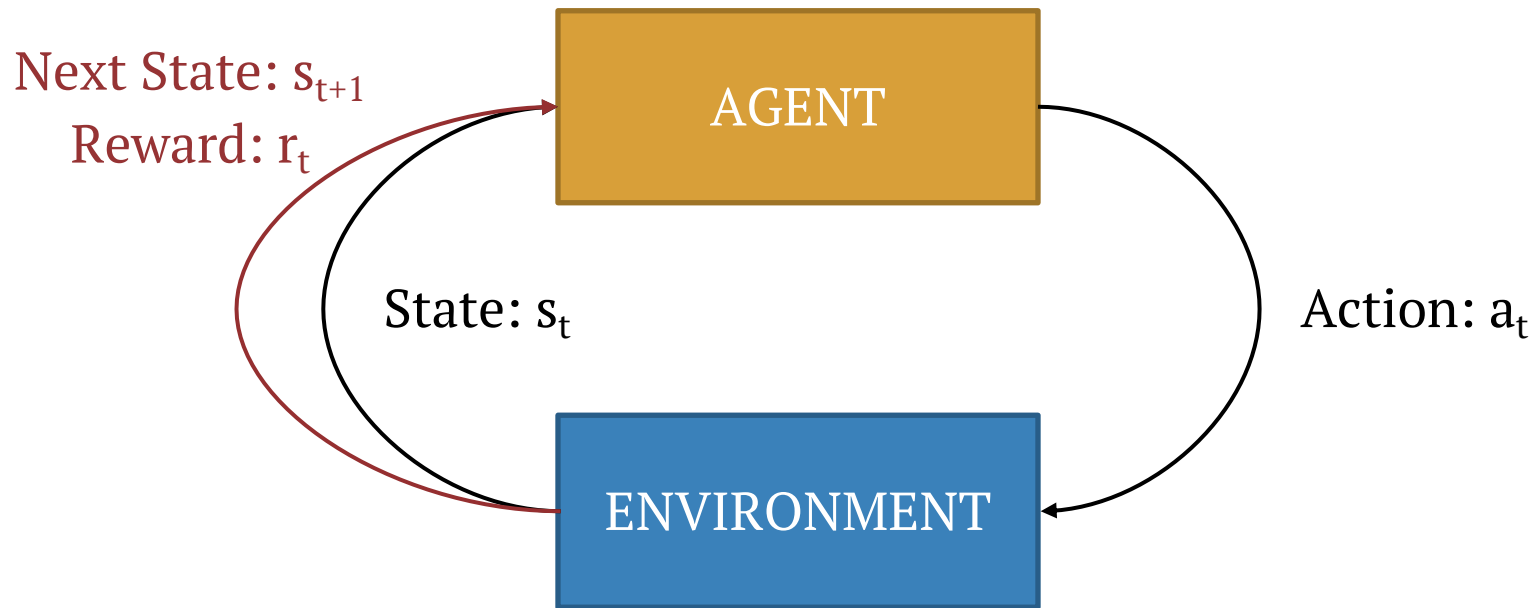
■ **State:**

■ **Action:**

■ **Reward:**



Reinforcement Learning



How can we mathematically formalize the RL problem?

Markov Decision Process

- ▣ Mathematical formulation of the RL problem –
- ▣ **Markov property:** Current state completely characterizes the state of the world

Defined by: $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{P}, \gamma)$

\mathcal{S} : set of possible states


\mathcal{A} : set of possible actions

\mathcal{R} : distribution of reward given (state, action) pair

\mathbb{P} : transition probability i.e. distribution over next state given (state, action) pair

γ : discount factor

Markov Decision Process: Example



Markov Decision Process

- At time step $t=0$, environment samples initial state $s_0 \sim p(s_0)$
- Then, for $t=0$ until done:
 - Agent selects action at a_t
 - Environment samples reward
$$r_t \sim R(.|s_t, a_t)$$
 - Environment samples next state
$$s_{t+1} \sim P(.|s_t, a_t)$$
 - Agent receives reward r_t and next state s_{t+1}

Markov Decision Process

- ▣ A policy π is a function from S to A that specifies what action to take in each state
- ▣ **Objective:** find policy π^* that maximizes cumulative discounted reward:

$$\sum_{t=0} \gamma^t r_t$$

$$= r_t + \gamma^t r_{t+1} + \gamma^2 r_{t+2} + \dots$$

A Simple MDP: Grid World

actions = {

1. right →

2. left ←

3. up ↑

4. down ↓

}


states

★			
			★


Set a negative “reward”
for each transition
(e.g. $r = -1$)

Objective: reach one of terminal states (greyed out)
in least number of actions

A Simple MDP: Grid World



Random Policy



Optimal Policy

Objective: reach one of terminal states (greyed out)
in least number of actions

The Optimal Policy π^*

- ▣ **Objective:** find optimal policy π^* that maximizes the sum of rewards
- ▣ To handle the randomness (initial state, transition probability, etc.), we need to maximize the **expected sum of rewards!**

$$\pi^* = \arg \max_{\pi} \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t | \pi \right] \quad \text{with } s_0 \sim p(s_0), a_t \sim \pi(\cdot | s_t), s_{t+1} \sim p(\cdot | s_t, a_t)$$

Value Function

- ▣ Following a policy produces sample trajectories (or paths): $(s_0, a_0, r_0), (s_1, a_1, r_1), \dots$
- ▣ **How good is a state?**
 - The **value function** at state s , is the expected cumulative reward from following the policy from state s :

$$V^\pi(s) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid s_0 = s, \pi \right]$$

Q-Value Function

- ▣ Following a policy produces sample trajectories (or paths): $(s_0, a_0, r_0), (s_1, a_1, r_1), \dots$
- ▣ **How good is a state-action pair?**
 - The **Q-value function** at state s and action a , is the expected cumulative reward from taking action a in state s and then following the policy:

$$Q^\pi(s, a) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid s_0 = s, a_0 = a, \pi \right]$$

Bellman Equation

- The optimal Q-value function Q^* is the maximum expected cumulative reward achievable from a given (state, action) pair:

$$Q^*(s, a) = \max_{\pi} \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid s_0 = s, a_0 = a, \pi \right]$$

- Q^* satisfies the following Bellman equation:

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

Solving for the Optimal Policy

- **Value iteration** algorithm: Use Bellman equation as an iterative update

$$Q_{i+1}(s, a) = \mathbb{E} \left[r + \gamma \max_{a'} Q_i(s', a') | s, a \right]$$

Q_i will converge to Q^* as $i \rightarrow \text{infinity}$

- **Problem:** Not scalable.
 - Must compute $Q(s, a)$ for every state-action pair.
 - If state, e.g. current game state pixels, is computationally infeasible to compute for entire state space!

Solving for the Optimal Policy: Q-Learning


- **Solution:** use a function approximator to estimate $Q(s,a)$ e.g. a neural network!
- **Q-learning** uses a function approximator to estimate the action-value function

$$Q(s, a; \theta) \approx Q^*(s, a)$$

 **function parameters (weights)**

- If the function approximator is a deep neural network => **Deep Q-learning!**

Simple Q-learning Algorithm Process



actions = {

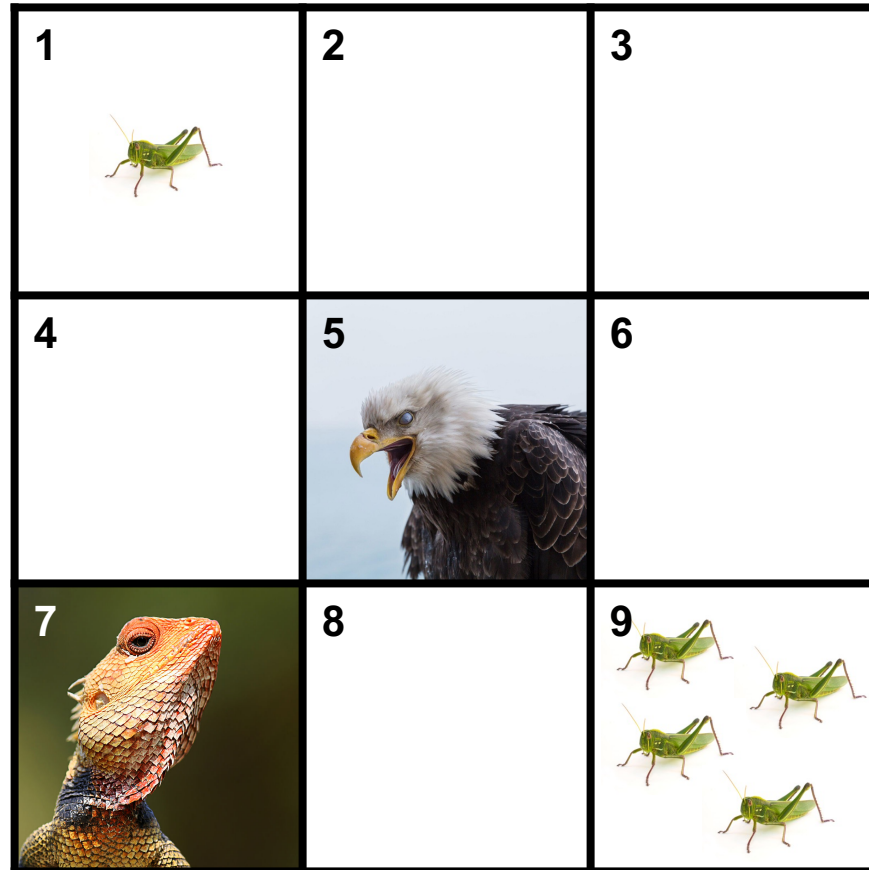
1. right →


2. left ←

3. up ↑

4. down ↓

}





Empty cell = -1

(1) Cricket = +1

(5) Eagle = -10 [End Game]

(9) Crickets = +10 [End Game]

Q-Table		Action			
		Left	Right	Up	Down
States	1 cricket	0	0	0	0
	2	0	0	0	0
	3	0	0	0	0
	4	0	0	0	0
	5 eagle	0	0	0	0
	6	0	0	0	0
	7 lizard	0	0	0	0
	8	0	0	0	0
	9 crickets	0	0	0	0

Q-Learning

- Goal: want to find a Q-function that satisfies the Bellman Equation:

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') | s, a \right]$$

- **Forward Feeding**

- Loss Function:

$$L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot)} \left[(y_i - Q(s, a; \theta_i))^2 \right]$$

- where

$$y_i = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) | s, a \right]$$

Q-Learning

- Goal: want to find a Q-function that satisfies the Bellman Equation:

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') | s, a \right]$$

- **Backpropagation**

- Gradient update (with respect to Q-function parameters θ):

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right] \nabla_{\theta_i} Q(s, a; \theta_i)$$

Update Q-Table

$$\underbrace{NewQ(s, a)} = \underbrace{Q(s, a)} + \underbrace{\alpha}_{\text{Learning Rate}} [\underbrace{R(s, a)}_{\text{Reward for taking that action at that state}} + \underbrace{\gamma}_{\text{Discount rate}} \underbrace{\max Q'(s', a')}_{\text{Maximum expected future reward given the new } s' \text{ and all possible actions at that new state}} - \underbrace{Q(s, a)}]$$

New Q value for that state and that action





Current Q value

Learning Rate

Reward for taking that action at that state

Discount rate

Maximum expected future reward **given the new s' and all possible actions at that new state**

1 	2	3
4	5 	6
7 	8	9 

Empty cell = -1

(1) Cricket = +1

(5) Eagle = -10 [End Game]

(9) Crickets = +10 [End Game]

Q-Table		Action			
		Left	Right	Up	Down
States	1 cricket	<i>UPDATE THE TABLE!!</i>			
	2				
	3				
	4				
	5 eagle				
	6				
	7 lizard				
	8				
	9 crickets				